

Part 2 — Lecture 5

TECH2: Introduction to Programming, Data, and Information Technology

Richard Foltyn

Norwegian School of Economics (NHH)

October 30, 2024

See GitHub repository for notebooks and data:

<https://github.com/richardfoltyn/TECH2-H24>

Contents

1	Plotting with Matplotlib	1
1.1	Line plots	1
1.2	Scatter plots	5
1.3	Plotting categorical data	7
1.4	Adding labels and annotations	8
1.5	Plot limits, ticks and tick labels	9
1.6	Adding straight lines	10
1.7	Object-oriented interface	11
1.8	Working with multiple plots (axes)	12
2	Plotting with pandas	15
2.1	Bar charts	15
2.2	Plotting time series data	17
2.3	Scatter plots	18
2.4	Box plots	20

1 Plotting with Matplotlib

In this unit, revisit the topic of plotting numerical data. Python itself does not have any built-in plotting capabilities, so we will be using `matplotlib` (MPL), the most popular graphics library for Python.

- For details on a particular plotting function, see the [official documentation](#).
- There is an official introductory [tutorial](#) which you can use along-side this unit.

In order to access the functions and objects from `matplotlib`, we first need to import them. The general convention is to use the namespace `plt` for this purpose:

```
import matplotlib.pyplot as plt
```


1.1 Line plots

One of the simplest plots we can generate using the `plot()` function is a line defined by a list of y -values.

```
[1]: # import matplotlib library
import matplotlib.pyplot as plt

# Plot list of integers
yvalues = [1, 4, 2, 3]
plt.plot(yvalues)
```


```
[1]: [<matplotlib.lines.Line2D at 0x7f3c335a9c40>]
```



We didn't even have to specify the corresponding x -values, as MPL automatically assumes them to be $[0, 1, 2, \dots]$. Usually, we want to plot for a given set of x -values like this:

```
[2]: # explicitly specify x-values
xvalues = [10, 20, 30, 40]
plt.plot(xvalues, yvalues)
```

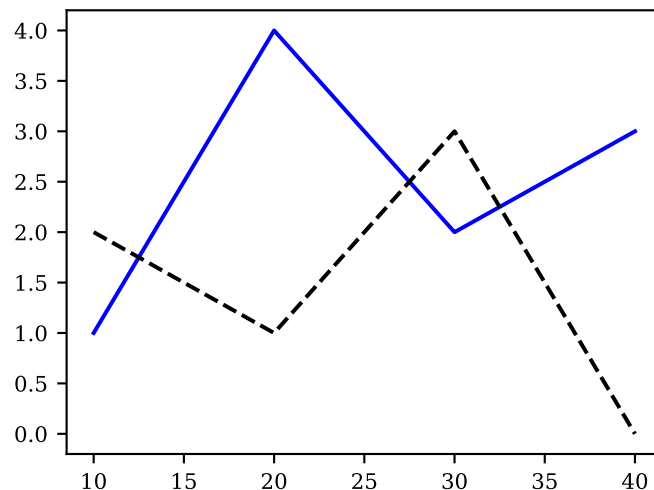
```
[2]: [<matplotlib.lines.Line2D at 0x7f3c31661a60>]
```



We can also specify multiple lines to be plotted in a single graph:

```
[3]: yvalues2 = [2.0, 1.0, 3.0, 0.0]
plt.plot(xvalues, yvalues, 'b-', xvalues, yvalues2, 'k--')
```

```
[3]: [<matplotlib.lines.Line2D at 0x7f3c316bf800>,
      <matplotlib.lines.Line2D at 0x7f3c31515160>]
```



The characters following each set of y -values are style specifications. The letters are short-hand notations for colors (see [here](#) for details):

- b: blue
- g: green
- r: red
- c: cyan
- m: magenta
- y: yellow
- k: black
- w: white

The remaining characters set the line styles. Valid values are

- - solid line
- -- dashed line
- - . dash-dotted line
- : dotted line

Additionally, we can append marker symbols to the style specification. The most frequently used ones are

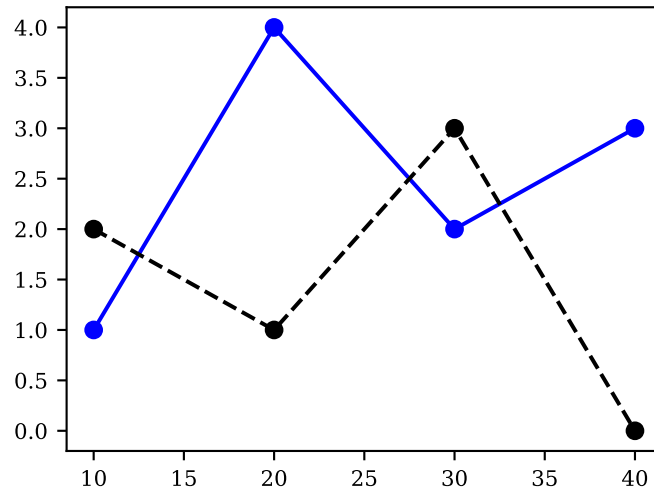
- o: circle
- s: square
- *: star
- x: x
- d: (thin) diamond

The whole list of supported symbols can be found [here](#).

Instead of passing multiple values to be plotted at once, we can also repeatedly call `plot()` to add additional elements to a graph. This is more flexible since we can pass additional arguments which are specific to one particular set of data, such as labels displayed in legends.

```
[4]: # Plot two lines by calling plot() twice
plt.plot(xvalues, yvalues, 'b-o')
plt.plot(xvalues, yvalues2, 'k--o')
```

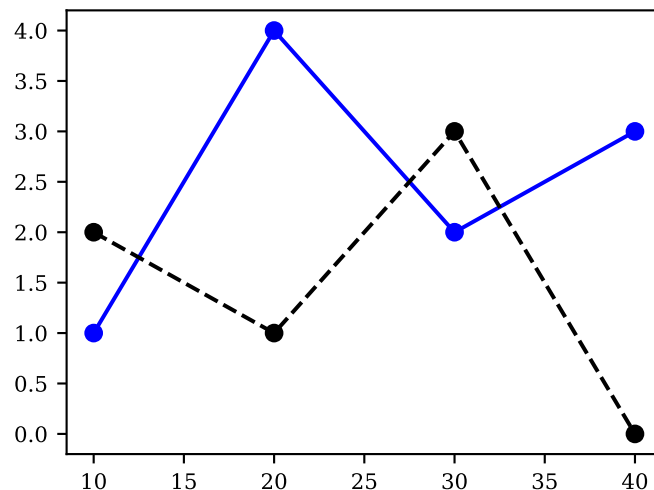
```
[4]: [matplotlib.lines.Line2D at 0x7f3c31570470]
```



Individual calls to `plot()` also allow us to specify styles more explicitly using keyword arguments:

```
[5]: # pass plot styles as explicit keyword arguments
plt.plot(xvalues, yvalues, color='blue', linestyle='-', marker='o')
plt.plot(xvalues, yvalues2, color='black', linestyle='--', marker='o')
```

```
[5]: [matplotlib.lines.Line2D at 0x7f3c3141a1b0]
```



Note that in the example above, we use named colors such as 'red' or 'blue' (see [here](#) for the complete list of named colors).

Matplotlib accepts abbreviations for the most common style definitions using the following shortcuts:

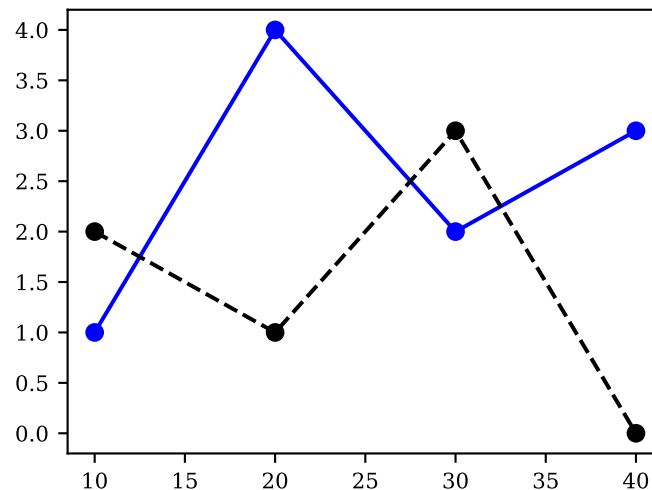
- `c` or `color`
- `ls` or `linestyle`
- `lw` or `linewidth`
- `ms` or `markersize`

See the section on *Other Parameters* in the `plot()` documentation for a complete list of arguments and their abbreviations.

We can write thus rewrite the above code as follows:

```
[6]: # abbreviate plot style keywords
plt.plot(xvalues, yvalues, c='blue', ls='-', marker='o')
plt.plot(xvalues, yvalues2, c='black', ls='--', marker='o')
```

```
[6]: [<matplotlib.lines.Line2D at 0x7f3c314966f0>]
```



Your turn. Use the data files located in the folder `../data/FRED` to perform the following tasks:

1. Load the data in `REALRATE.csv` and `FEDFUNDS.csv` into two different DataFrames. The files contain the 1-year real interest rate and the Federal Funds rate at monthly frequency.

Hint: Use `pd.read_csv(..., parse_dates=['DATE'], index_col='DATE')` to automatically parse strings stored in the `DATE` column as dates and set `DATE` as the index.

2. Plot the `REALRATE` using a blue dashed line with line width 0.5 and `FEDFUNDS` using an orange line with line width 0.75 in the same figure.

1.2 Scatter plots


We use the `scatter()` function to create scatter plots in a similar fashion to line plots:

```
[7]: import matplotlib.pyplot as plt
import numpy as np

# Create 50 uniformly-spaced values on unit interval
xvalues = np.linspace(0.0, 1.0, 50)
# Draw random numbers
yvalues = np.random.default_rng(123).random(50)

plt.scatter(xvalues, yvalues, color='blue')
```

```
[7]: <matplotlib.collections.PathCollection at 0x7f3c315b8e00>
```




We could in principle create scatter plots using `plot()` by turning off the connecting lines. However, `scatter()` allows us to specify the color and marker size as collections, so we can vary these for every point. `plot()`, on the other hand, imposes the same style on all points plotted in that particular function call.

```
[8]: # Draw random marker sizes
size = np.random.default_rng(456).random(len(yvalues)) * 150.0

# plot with point-specific marker sizes
plt.scatter(xvalues, yvalues, s=size)
```

```
[8]: <matplotlib.collections.PathCollection at 0x7f3c3134dcd0>
```



Your turn. Use the data files located in the folder `../data/FRED` to perform the following tasks:

1. Load the data in `REALRATE.csv` and `FEDFUNDS.csv` into two different DataFrames. The files contain the 1-year real interest rate and the Federal Funds rate at monthly frequency.

Hint: Use `pd.read_csv(..., parse_dates=['DATE'], index_col='DATE')` to automatically parse strings stored in the `DATE` column as dates and set `DATE` as the index.

2. Create a scatter plot of `REALRATE` (on the y-axis) against `FEDFUNDS` (on the x-axis). Specify the arguments `edgecolors='blue'` and `color='none'` to plot the data as blue rings.

Hint: `scatter()` requires the arguments to be of equal length, which is not the case here. Restrict the data to the same time period, for example by using `join(..., how='inner')`.

1.3 Plotting categorical data

Instead of numerical values on the x -axis, we can also plot categorical variables using the function `bar()`.

For example, assume we have four categories and each has an associated numerical value:

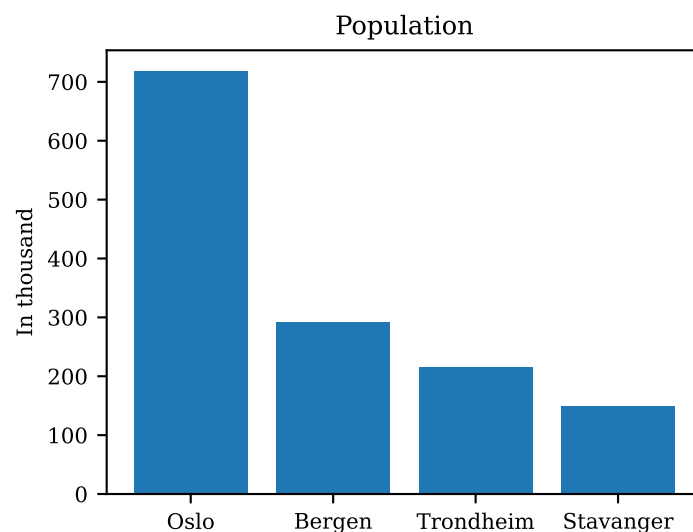
```
[9]: import matplotlib.pyplot as plt

# Define category labels
municipality = ['Oslo', 'Bergen', 'Trondheim', 'Stavanger']
# Population in thousand
population = np.array([717710, 291940, 214565, 149048]) / 1000

# Create bar chart
plt.bar(municipality, population)

# Add overall title
plt.title('Population')
plt.ylabel('In thousand')
```

```
[9]: Text(0, 0.5, 'In thousand')
```




We use `barh()` to create *horizontal* bars:

```
[10]: plt.barh(municipality, population)
plt.title('Population')
```

```
plt.xlabel('In thousand')
```

```
[10]: Text(0.5, 0, 'In thousand')
```



1.4 Adding labels and annotations

Matplotlib has numerous functions to add labels and annotations:

- Use `title()` and `suptitle()` to add titles. The latter adds a title for the whole figure, which might span multiple plots (axes).
- We can add axis labels by calling `xlabel()` and `ylabel()`.
- To add a legend, call `legend()`, which in its most simple form takes a list of labels which are in the same order as the plotted data.
- Use `text()` to add additional text at arbitrary locations.


```
[11]: import matplotlib.pyplot as plt

xvalues = [0, 1, 2, 3]
yvalues = [1, 4, 2, 3]
yvalues2 = [2.0, 1.0, 3.0, 0.0]

plt.plot(xvalues, yvalues, 'r', xvalues, yvalues2, 'b')
plt.suptitle('Figure title')
plt.title('Axes-specific title')
plt.xlabel('Label for x-axis')
plt.ylabel('Label for y-axis')
plt.legend(['Red line', 'Blue line'])

# Adds text at data coordinates (0.05, 0.05)
plt.text(0.05, 0.05, 'More text')
```

```
[11]: Text(0.05, 0.05, 'More text')
```

1.5 Plot limits, ticks and tick labels

We adjust the plot limits, ticks and tick labels as follows:

- Plotting limits are set using the `xlim()` and `ylim()` functions. Each accepts a tuple (min,max) to set the desired range.
- Ticks and tick labels can be set by calling `xticks()` or `yticks()`.


```
[12]: import matplotlib.pyplot as plt
import numpy as np

xvalues = np.linspace(0.0, 2*np.pi, 50)
plt.plot(xvalues, np.sin(xvalues))

# Set major ticks for x and y axes, and xtick labels.
# We can use LaTeX code in labels!
plt.xticks([0.0, np.pi, 2*np.pi], ['0', r'$\pi$', r'$2\pi$'])
plt.yticks([-1.0, 0.0, 1.0])

# Adjust plot limits in x and y direction
plt.xlim((-0.5, 2*np.pi + 0.5))
plt.ylim((-1.1, 1.1))
```

```
[12]: (-1.1, 1.1)
```



1.6 Adding straight lines

Quite often, we want to add horizontal or vertical lines to highlight a particular value. We can do this using the functions

- `axhline()` to add a *horizontal* line at a given *y*-value.
- `axvline()` to add a *vertical* line at a given *x*-value.
- `axline()` to add a line defined by two points or by a single point and a slope.

Example: Adding horizontal and vertical lines

Consider the sine function from above. We can add a horizontal line at 0 and two vertical lines at the points where the function attains its minimum and maximum as follows:


```
[13]: import matplotlib.pyplot as plt
import numpy as np

# Plot sine function (same as above)
xvalues = np.linspace(0.0, 2*np.pi, 50)
plt.plot(xvalues, np.sin(xvalues))
plt.xticks([0.0, np.pi, 2*np.pi], ['0', r'$\pi$', r'$2\pi$'])
plt.yticks([-1, 0, 1])

# Add black dashed horizontal line at y-value 0
plt.axhline(0.0, lw=0.5, ls='--', c='black')

# Add red dashed vertical lines at maximum / minimum points
plt.axvline(0.5*np.pi, lw=0.5, ls='--', c='red')
plt.axvline(1.5*np.pi, lw=0.5, ls='--', c='red')
```

```
[13]: <matplotlib.lines.Line2D at 0x7f3c31168bf0>
```



1.7 Object-oriented interface

So far, we have only used the so-called pyplot interface which involves calling *global* plotting functions from `matplotlib.pyplot`. This interface is intended to be similar to Matlab, but is also somewhat limited and less clean.

We can instead use the object-oriented interface (called this way because we call methods of the [Figure](#) and [Axes](#) objects). While there is not much point in using the object-oriented interface in a Jupyter notebook when we want to create a single graph, it should be the preferred method when writing re-usable code in Python files.

To use the object-oriented interface, we need to get figure and axes objects. The easiest way to accomplish this is using the `subplots()` function, like this:

```
fig, ax = plt.subplots()
```

We then use methods of the Axes object returned by `subplots()` instead of the functions we have used so far. For example, instead of `plot()`, we use the `Axes.plot()` method.


As an example, we recreate the graph from the section on labels and annotations using the object-oriented interface:

```
[14]: import matplotlib.pyplot as plt

xvalues = [0, 1, 2, 3]
yvalues = [1, 4, 2, 3]
yvalues2 = [2.0, 1.0, 3.0, 0.0]

fig, ax = plt.subplots()
ax.plot(xvalues, yvalues, color='red', label='Red line')
ax.plot(xvalues, yvalues2, color='blue', label='Blue line')
ax.set_xlabel('Label for x-axis')
ax.set_ylabel('Label for y-axis')
ax.legend()
ax.set_title('Plot with object-oriented interface')
ax.text(0.05, 0.05, 'More text')
```

```
[14]: Text(0.05, 0.05, 'More text')
```



The code is quite similar to the previous section, except that attributes are set using the `set_xxx()` methods of the `ax` object. For example, instead of calling `xlim()`, we use `ax.set_xlim()`.


1.8 Working with multiple plots (axes)

The object-oriented interface becomes particularly useful if we want to create multiple axes (or figures). This can also be achieved using the pyplot programming model but is somewhat more obscure.

For example, to create a row with two subplots, we use:

```
[15]: import matplotlib.pyplot as plt

# Create one figure with 2 axes objects, arranged as two columns in a single row
fig, ax = plt.subplots(1, 2, figsize=(4.5, 2.0))
```




With multiple axes objects in a single figure (as in the above example), the `ax` returned by `subplots()` is a NumPy array. Its elements map to the individual panels within the figure in a natural way.

We can visualise this mapping for the case of a single row and two columns as follows:

```
[16]: fig, ax = plt.subplots(1, 2, figsize=(3.5, 1.5))
for i, axes in enumerate(ax):
    # Turn off ticks of both axes
    axes.set_xticks(())
    axes.set_yticks(())
    # Label axes object
    text = f'ax[{i}]'
```


```
axes.text(0.5, 0.5, text, transform=axes.transAxes, va='center',
         ha='center', fontsize=18)
```



Don't worry about the details of how this graph is generated, the only take-away here is how axes objects are mapped to the panels in the figure.

If we request panels in two dimensions, the ax object will be a 2-dimensional array, and the mapping of axes objects to panels will look like this instead:

```
[17]: # Request figure with 2 rows, 3 columns
fig, ax = plt.subplots(2, 3, figsize=(6,4))
for i, axrow in enumerate(ax):
    for j, axes in enumerate(axrow):
        # Turn off ticks of both axes
        axes.set_xticks(())
        axes.set_yticks(())
        # Label axes object
        text = f'ax[{i},{j}]'
        axes.text(0.5, 0.5, text, transform=axes.transAxes, va='center',
                 ha='center', fontsize=18)
```



Example: Create a plot with 2 panels

We can use the elements of ax to plot into individual panels:


```
[18]: import matplotlib.pyplot as plt
import numpy as np

fig, ax = plt.subplots(1, 2, sharex=True, sharey=True, figsize=(4.5, 2.0))
xvalues = np.linspace(-1.0, 1.0, 50)
```

```
# Plot into first column
ax[0].plot(xvalues, xvalues)

# Plot into second column
ax[1].plot(xvalues, 2*xvalues**2.0 - 1)
```

[18]: [<matplotlib.lines.Line2D at 0x7f3c30dde8a0>]




Example: Create a figure with 2 rows and 2 columns

```
[19]: # create figure with 2 rows, 2 columns
fig, ax = plt.subplots(2, 2, sharex=True, sharey=True)

xvalues = np.linspace(-1.0, 1.0, 50)

# Plot the first four powers of x
exponent = 1
for i in range(2):
    for j in range(2):
        yvalues = xvalues**exponent
        ax[i,j].plot(xvalues, yvalues)
        # Increment exponent for next subplot
        exponent += 1
```



Note the use of `sharex=True` and `sharey=True`. This tells matplotlib that all axes share the same plot limits, so the tick labels can be omitted in the figure's interior to preserve space.

Your turn. Create a figure with 3 columns (on a single row) and plot the following functions on the interval $[0, 6]$:

- Subplot 1: $y = \sin(x)$
- Subplot 2: $y = \sin(2 * x)$
- Subplot 3: $y = \sin(4 * x)$

Hint: The sine function can be imported from NumPy as `np.sin()`.

2 Plotting with pandas

Pandas does not implement its own graphics library, but provides convenient wrappers around Matplotlib functions that can be used to quickly visualize data stored in DataFrames. Alternatively, we can extract the numerical data and pass it to Matplotlib's routines manually.

2.1 Bar charts

Let's return to our municipality population data. To plot population numbers as a bar chart, we can directly use pandas's `plot.bar()`:

```
[20]: import pandas as pd


# Path to local data/ folder
DATA_PATH = '../data'

# Read in population data for Norwegian municipalities
df = pd.read_csv(f'{DATA_PATH}/population_norway.csv')

# Keep only the first 15 observations
df = df.iloc[:15]

# Create bar chart, specify figure size, rotate x-ticklabels by 45 degrees
df.plot.bar(x='Municipality', y='Population', rot=45, figsize=(6,3))
```

```
[20]: <Axes: xlabel='Municipality'>
```



Alternatively, we can construct the graph ourselves using Matplotlib:

```
[21]: import matplotlib.pyplot as plt

# Extract municipality names
labels = df['Municipality']


# Extract population numbers
values = df['Population']

# Create new figure with desired size
plt.figure(figsize=(6, 3))

# Create bar chart
plt.bar(labels, values)

# Add legend
plt.legend(['Population'])

# Rotate ticks
plt.tick_params(axis='x', labelrotation=45)
```

Matplotlib's functions usually directly work with pandas's data structures, In cases where they don't, we can convert a DataFrame or Series object to a NumPy array using the `to_numpy()` method.

2.2 Plotting time series data

To plot time series data, we can use the `DataFrame.plot()` method which optionally accepts arguments to specify which columns should be used for the x -axis and which for the y -axis. We illustrate this using the US unemployment rate at annual frequency.


```
[22]: import numpy as np
import pandas as pd

# Path to FRED.csv; DATA_PATH variable was defined above!
filepath = f'{DATA_PATH}/FRED.csv'

# Read CSV data
df = pd.read_csv(filepath, sep=',')

# Plot unemployment rate by year
df.plot(x='Year', y='UNRATE', ylabel='Unemployment rate (%)')
```

```
[22]: <Axes: xlabel='Year', ylabel='Unemployment rate (%)'>
```



Your turn. Use the data files located in the folder `../data` to perform the following tasks:

1. Load the macroeconomic time series data from `FRED.csv`.

Hint: Use `pd.read_csv(..., parse_dates=['DATE'])` to automatically set parse strings as dates.

2. Create a line plot, showing both the unemployment rate `UNRATE` and the inflation rate `INFLATION` in a single graph.

2.3 Scatter plots

Using the `DataFrame.plot.scatter()` method, we can generate scatter plots, plotting one column against another. To illustrate, we plot the US unemployment rate against inflation in any given year over the post-war period.

Note that you can pass additional arguments to pandas's version of `scatter()` which are passed on to Matplotlib's `scatter()`, for example `edgecolors`.


```
[23]: # Path to FRED.csv; DATA_PATH variable was defined above!
      filepath = f'{DATA_PATH}/FRED.csv'

      # Read CSV data
      df = pd.read_csv(filepath, sep=',')

      # Compute annual inflation as growth rate of CPI
      df['Inflation'] = df['CPI'].pct_change() * 100.0

      df.plot.scatter(
          x='UNRATE',
          y='Inflation',
          color='none',
          edgecolors='steelblue',
          xlabel='Unemployment rate (%)',
          ylabel='Inflation (%)'
      )
```

```
[23]: <Axes: xlabel='Unemployment rate (%)', ylabel='Inflation (%)'>
```




Pandas also offers the convenience function `scatter_matrix()` which lets us easily create pairwise scatter plots for more than two variables:

```
[24]: from pandas.plotting import scatter_matrix

# Continue with DataFrame from previous example, compute GDP growth
df['GDP_growth'] = df['GDP'].pct_change() * 100.0

# Columns to include in plot
columns = ['GDP_growth', 'Inflation', 'UNRATE']

# Use argument diagonal='kde' to plot kernel density estimate
# in diagonal panels
axes = scatter_matrix(
    df[columns],
    figsize=(6, 6),
    diagonal='kde',          # plot kernel density along diagonal
    s=70,                   # marker size
    color='none',
    edgecolors='steelblue',
    alpha=1.0,
)
```



2.4 Box plots

To quickly plot some descriptive statistics, we can use the `DataFrame.plot.box()` provided by pandas. This plot shows the median, the interquartile range (25th to 75th percentile) and the outliers of some underlying data.

We demonstrate this by plotting the distribution of post-ware GDP growth, inflation and the unemployment rate in the US:

```
[25]: import numpy as np
import pandas as pd

# Path to FRED.csv; DATA_PATH variable was defined above!
filepath = f'{DATA_PATH}/FRED.csv'

# Read CSV data
df = pd.read_csv(filepath, sep=',')

# Compute annual growth rates (in percent)
df['GDP_growth'] = df['GDP'].pct_change() * 100.0
df['Inflation'] = df['CPI'].pct_change() * 100.0

# Include only the following columns in plot
columns = ['GDP_growth', 'Inflation', 'UNRATE']

# Create box plot. Alternatively, use df.plot(kind='box')
```

```
df[columns].plot.box(ylabel='Relative change in % (annual)')
```

```
[25]: <Axes: ylabel='Relative change in % (annual)'\>
```

